

National workshop on  
Applications of Deep Learning in Computer Vision  
SSN Engineering College

# Introduction to Deep Learning

S Milton Rajendram  
Computer Science and Engineering  
SSN College of Engineering  
[miltonrs@ssn.edu.in](mailto:miltonrs@ssn.edu.in)



10 February 2022

# Logistic Regression

# Linear classification

- ▶  $\mathcal{D} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$  ▶  $y$  is binary-valued (discrete-valued)

- ▶ Linear combination of features of  $\mathbf{x}$

$$w_0 + w_1x_1 + \dots + w_dx_d = w_0 + \sum_{i=1}^d w_ix_i$$

$$w_0x_0 + w_1x_1 + \dots + w_dx_d = \sum_{i=0}^d w_ix_i, x_0 = 1$$

- ▶ Vectors in matrix notation


$$\mathbf{x} = \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_d \end{bmatrix} \quad \mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ x_d \end{bmatrix}$$

$$\hat{y} = \begin{cases} +1 & \text{if } \mathbf{w}^T \mathbf{x} > 0 \\ -1 & \text{if } \mathbf{w}^T \mathbf{x} < 0 \end{cases}$$

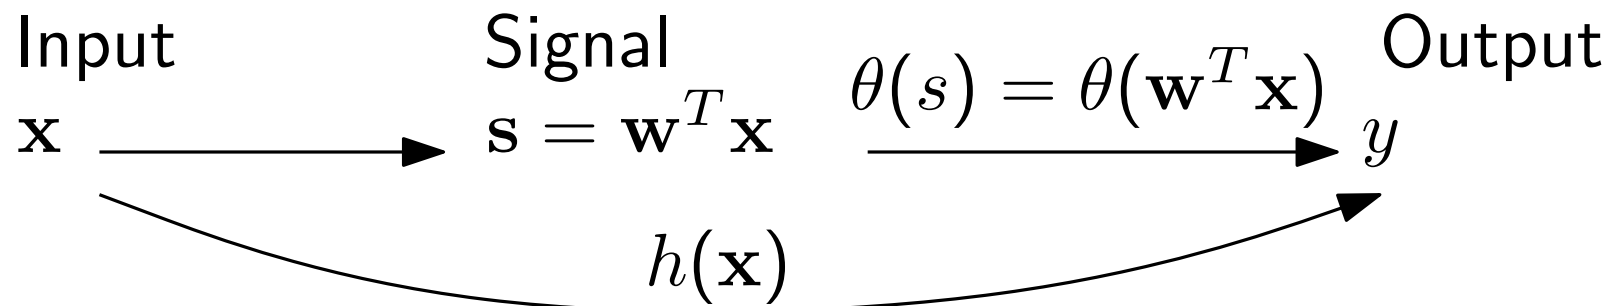
$$\hat{y} = \text{sign}(\mathbf{w}^T \mathbf{x})$$

- ▶ Linear combination


$$\sum_{i=0}^d w_ix_i = \mathbf{w}^T \mathbf{x} = \mathbf{w} \cdot \mathbf{x}$$




# Linear models



Linear classification  
 $h(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x})$




Linear regression  
 $h(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$



# Logistic regression


Linear classification

$$h(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x})$$




Logistic regression

$$0 \leq h(\mathbf{x}) = \theta(\mathbf{w}^T \mathbf{x}) \leq 1$$



Linear regression

$$h(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$$




# Logistic regression — overview


- ▶  $h(\mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}} = P(y|\mathbf{x})$ , hypothesis set
- ▶  $e(h(\mathbf{x}_i), y) = \ln(1 + e^{-y_i \mathbf{w}^T \mathbf{x}_i})$ , cross-entropy error
- ▶  $E(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N \ln(1 + e^{-y_i \mathbf{w}^T \mathbf{x}_i})$ , in-sample error or likelihood
- ▶ Minimize in-sample error  $E(\mathbf{w})$
- ▶ Gradient  $\nabla_{\mathbf{w}} E(\mathbf{w})$ , in which direction of  $\mathbf{w}$  error decreases most rapidly and how much
- ▶ Gradient descent:  
$$\mathbf{w}_{(t+1)} \leftarrow \mathbf{w}_{(t)} + \eta \nabla_{\mathbf{w}} E(\mathbf{w})$$

# Linear is limited

Data




Hypothesis



- ▶ Credit limit is affected by 'years in residence'
- ▶ But **not** in a linear way!

# Idea

- ▶ Generalization of perceptron.
- ▶ Can model complex (non-linear) target function.
- ▶ Fits the data by minimizing in-sample error  $E$
- ▶ Flexible, but risk of overfitting.






# Neural Networks

# Idea

- ▶ Generalization of perceptron.
- ▶ Can model complex (non-linear) target function.
- ▶ Fits the data by minimizing in-sample error  $E$
- ▶ Flexible, but risk of overfitting.

# Multi-layer perceptron (MLP) of XNOR



3-layer feed-forward


$$\begin{array}{ll}
 h_1(\mathbf{x}) = \mathbf{w}_1^T \mathbf{x} & h_1(\mathbf{x}).\overline{h_2(\mathbf{x})} \\
 h_2(\mathbf{x}) = \mathbf{w}_2^T \mathbf{x} & \overline{h_1(\mathbf{x})}.h_2(\mathbf{x})
 \end{array}
 \quad
 \frac{h_1(\mathbf{x}).\overline{h_2(\mathbf{x})}}{h_1(\mathbf{x}).h_2(\mathbf{x})} +$$

- Between input and output, multiple **hidden** layers.


# Feed-forward MLP

- ▶ Each layer feed-forward to the next layer only. Layer  $i$  to layer  $i + 1$  only.
- ▶ No feed-backward, no jump forward to other layers.
- ▶ Perceptron has just input and output.
- ▶ Multi-layer perceptron
  - ▶ input layer with  $d_{in}$  nodes,
  - ▶ output layer with  $d_{out}$  node,
  - ▶ 2 **hidden layers** with 3 nodes (**hidden units**) each.
- ▶ **Architecture** of MLP: number of hidden layers and number of hidden units in each layer.
- ▶ A 3-layer MLP with suitably many hidden units can model any target function — can fit any data.
- ▶ May **overfit**! May not generalize.

# Neural network



# Activation function



- ▶ Desirable that activation functions are differentiable, and
- ▶ Expressed in terms of the function itself.

▶ Tanh: output  $\theta(s)$   $\theta(s) = \tanh(s) = \frac{e^s - e^{-s}}{e^s + e^{-s}}$

▶ ReLU: output  $\theta(s) = \max(s, 0)$

▶ Sigmoid: output  $\theta(s) = \frac{1}{1 + e^{-s}}$

# Output-input for one layer

$$\begin{array}{ccc}
 \mathbf{x}^{(k-1)} \xrightarrow{W^{(k)}} & & \mathbf{s}^{(k)} \xrightarrow{\theta} \mathbf{x}^{(k)} \\
 \begin{bmatrix} x_0^{(k-1)} \\ x_1^{(k-1)} \\ \vdots \\ x_R^{(k-1)} \end{bmatrix} & \begin{bmatrix} w_{01} & w_{02} & \dots & w_{0S} \\ w_{11} & w_{12} & \dots & w_{1S} \\ w_{21} & w_{22} & \dots & w_{2S} \\ \vdots & \vdots & & \vdots \\ w_{R1} & w_{R2} & \dots & w_{RS} \end{bmatrix} & \begin{bmatrix} s_0^{(k)} \\ s_1^{(k)} \\ \vdots \\ s_S^{(k)} \end{bmatrix} \quad \begin{bmatrix} x_0^{(k)} \\ x_1^{(k)} \\ \vdots \\ x_S^{(k)} \end{bmatrix} \\
 R \times 1 & R \times S & S \times 1 \quad S \times 1
 \end{array}$$

$$(\mathbf{s}^{(k)})^T = (\mathbf{x}^{(k-1)})^T \cdot \mathbf{W}^{(k)}$$

$$\mathbf{x}^{(k)} = \theta(\mathbf{s}^{(k)})$$

# Forward propagation idea

$$\mathbf{x}^{(0)} \xrightarrow{W^{(1)}} \mathbf{s}^{(1)} \xrightarrow{\theta} \mathbf{x}^{(1)} \xrightarrow{W^{(2)}} \mathbf{y}$$

$$\begin{bmatrix} x_0^{(0)} \\ x_1^{(0)} \\ \vdots \\ x_{D_1}^{(0)} \end{bmatrix} \quad \begin{bmatrix} s_0^{(1)} \\ s_1^{(1)} \\ \vdots \\ s_H^{(1)} \end{bmatrix} \quad \begin{bmatrix} x_0^{(1)} \\ x_1^{(1)} \\ \vdots \\ x_H^{(1)} \end{bmatrix} \quad \begin{bmatrix} y_0^{(2)} \\ y_1^{(2)} \\ \vdots \\ y_{D_2}^{(2)} \end{bmatrix}$$

$$\mathbf{s} = \mathbf{x} \cdot \mathbf{W}^{(1)}$$

$$\mathbf{h} = \theta(\mathbf{s})$$

$$\mathbf{y} = \mathbf{h} \cdot \mathbf{W}^{(2)}$$

$$\mathbf{x}_{(1 \times D_1)} \xrightarrow{W1_{(D_1 \times H)}} \mathbf{s}_{(1 \times H)} \xrightarrow{\theta} \mathbf{h}_{(1 \times H)} \xrightarrow{W2_{(H \times D_2)}} \mathbf{y}_{(1 \times D_2)}$$

$$\begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_{D_1} \end{bmatrix} \quad \begin{bmatrix} s_0 \\ s_1 \\ \vdots \\ s_H \end{bmatrix} \quad \begin{bmatrix} h_0 \\ h_1 \\ \vdots \\ h_H \end{bmatrix} \quad \begin{bmatrix} y_0 \\ y_1 \\ \vdots \\ y_{D_2} \end{bmatrix}$$

$$D_1 \times 1$$

$$H \times 1$$

$$H \times 1$$

$$D_2 \times 1$$



# Training error


- ▶ Training error

$$\begin{aligned} E(\mathbf{w}) &= \frac{1}{N} \sum_{i=1}^N e_i \\ &= \frac{1}{N} \sum_{i=1}^N (\hat{y}_i - y_i)^2 \end{aligned}$$

- ▶ To minimize  $E(\mathbf{w})$ , solve  $\nabla E(\mathbf{w}) = 0$  for  $\mathbf{w}$ .

# Iterative method: gradient descent

- ▶ General method for non-linear optimization
- ▶ Start at  $\mathbf{w}_{(0)}$ , descend the surface along the steepest slope
- ▶ Let  $\hat{\mathbf{v}}$  (unit vector) be the direction of steepest slope.
- ▶ Take fixed-size steps along  $\hat{\mathbf{v}}$ :  $\mathbf{w}_{(t+1)} = \mathbf{w}_{(t)} + \eta \hat{\mathbf{v}}$



# Stochastic Gradient Descent (SGD)

1. Initialize  $\mathbf{w}_{(0)}$
2. for  $t = 1, 2, 3, \dots$  until  $\mathbf{w}$  converges do
3.     for  $i = 1 \dots N$  do
4.         Compute gradient  
                     $\nabla E(\mathbf{w})$
5.         Update weight:  $\mathbf{w}(t + 1) \leftarrow \mathbf{w}(t) - \eta \nabla E(\mathbf{w}(t))$
6. Return  $\mathbf{w}$

- ▶ For **updating  $\mathbf{w}$  once**, it considers **only one** training example, chosen at random.
- ▶ Converges faster, but may oscillate.

# Backpropagation idea

$$\text{Sensitivity of layer } k, \delta^{(k)} = \frac{\partial e}{\partial s^{(k)}}$$

- ▶ Partial derivatives in layer  $k - 1$  in terms of partial derivatives of layer  $k$


$$\frac{\partial e(\mathbf{w})}{\partial W^{(k)}} = x^{(k-1)} \times (\delta^{(k)})^T$$

$$\delta^{(k)} = \theta'(s^{(k)}) \otimes [W^{(k+1)} \times \delta^{(k+1)}]_1^{D_k}$$

$$\delta^{(L)} = 2(\hat{\mathbf{y}} - \mathbf{y})\theta'(\hat{\mathbf{y}})$$

# Backpropagation algorithm

- 1: Initialize all weights  $w_{ij}^k$  at random
- 2: for  $t = 1, 2 \dots$  do
- 4:   Forward: Compute all  $x_j^{(k)}$
- 5:   Backward: Compute all  $\delta_j^{(k)}$
- 6:   Update weights:  $w_{ij}^{(k)} \leftarrow w_{ij}^{(k)} - \eta x_i^{(k-1)} \delta_j^{(k)}$
- 7: Return the final weights  $w_{ij}^k$



# Issues in Back-propagation algorithm

- ▶ It has been shown that finding weights to minimize error is NP-complete [Blum and Rivest, 1989]
- ▶ Not guaranteed to converge
- ▶ Over-fitting to the training samples
- ▶ Error minimization process can get trapped in a local minima
- ▶ Extremely slow convergence
- ▶ Selection of initial weights may be critical
- ▶ Selection of optimal parameters
- ▶ Extremely difficult to explain the model and its predictions

# Deep Neural Networks

# Large datasets

- ▶ Today, we have very large datasets available – especially for images, videos, speech, and text.
- ▶ However, traditional **feature engineering** and classifiers based on given-representations have saturated – “shallow” networks have relatively small model parameters to capture information content of large data sets.
- ▶ More than two hidden layers.
- ▶ Very deep neural networks: up to hundreds of layers.



# Feature Engineering

- ▶ Major focus of machine learning has been selecting and extracting appropriate features
- ▶ Images: color, texture, edges, connected components, shapes, moments, SIFT, SURF, HOG, Wavelets-based, etc.
- ▶ Text: n-grams, character n-grams, POS tags, Named Entities, tags from shallow parsing, word sense disambiguation, etc.
- ▶ Speech: Mel Cepstral coefficients (MFCC), energy, zero crossing rates, pitch, timbre, etc.
- ▶ Which features are suitable for a given task? Leads to feature engineering.
- ▶ Which classifier is better? Ensemble of classifiers?
- ▶ Is a machine learning algorithm really “intelligent”?

# Challenges with deep neural networks

- ▶ Vanishing gradients
- ▶ Extremely slow convergence
- ▶ Lack of generalization
- ▶ Lack of large data – number of model parameters increases with depth
- ▶ Lack of labeled data – supervised learning requires examples with corresponding targets
- ▶ Computational inefficiency

# Vanishing gradient in backpropagation

- ▶ Network parameters are updated proportional to the partial derivative of the cost function w.r.t. the current parameters in each iteration of training.
- ▶  $\tanh$  has gradients in the range  $(0, 1)$
- ▶ Chain rule multiplies  $n$  of these small numbers to compute gradients of the leftmost layers in an  $n$ -layer network.
- ▶ Gradient decreases exponentially with  $n$ .
- ▶ Effectively preventing some parameters from changing their values.
- ▶ Earlier layers train very slowly, if at all.
- ▶ May completely stop the neural network from further training.
- ▶ Several improvements combined together: ReLU, LSTM, skip connections used in residual neural networks.

# New deep learning ideas

- ▶ New activation functions
- ▶ Regularization methods
- ▶ Initialization methods
- ▶ Data augmentation
- ▶ Optimization techniques
- ▶ GPU-based and distributed algorithms

# Activation Function: Rectifier

- ▶ Rectifier is a simple activation function defined as

$$f(x) = \max(0, x)$$

- ▶ Neuron using rectifier is popularly known as ReLU (Rectified Linear Unit)
- ▶ One variation of this is “softplus” function

$$f(x) = \ln(1 + e^x)$$

$$f'(x) = \frac{e^x}{1 + e^x} = \frac{1}{1 + e^{-x}}$$

- ▶ ReLU converges faster than networks using traditional 'tanh' [Krizhevsky et al., 2012]
- ▶ Variations of ReLU

# Activation Function: Softmax

- ▶ Squashes a  $k$ -dimension vector of real values to  $k$ -dimension vector of real values in the range  $(0,1)$  that add up to 1
- ▶ Converts a  $k$ -dimension vector to probability distribution over  $k$  different possible outcomes
- ▶ Useful in multi-class classification
- ▶ For example, softmax of  $[1,2,3,4,1,2,3]$  is  $[0.024, 0.064, 0.175, 0.475, 0.024, 0.064, 0.175]$

# Regularization methods: Dropout

- ▶ Set the output of a hidden neuron to 0 with a probability of 0.5 (effectively “drop” those neurons out).
- ▶ Dropped out neurons do not take part in forward pass and back-propagation.
- ▶ Hence, different architecture is sampled every time, but these networks share weights.
- ▶ Reduces the complex co-adaption of neurons and reduces the model parameters to be learned.
- ▶ However, dropout almost doubles the number of iterations required.

# Data augmentation

- ▶ We may not have enough data for learning a large number of model parameters (especially labeled data)
- ▶ Augment the existing data with some “label-preserving” operations
  - ▶ Image translations and horizontal reflections. Given, 256x256 images, extract random 224x224 patches and their horizontal reflections – object labels do not change!
  - ▶ Alter the intensities of RGB channels – one idea may be to perform PCA to find principal components and add a magnitude of this to the intensities.




# Stochastic Gradient Descent

- ▶ Stochastic Gradient Descent (SGD) incrementally updates the weights for each training example.
- ▶ This may result in wide swings the weights (which may be good sometimes!)
- ▶ To overcome this problem, learning rate is slowly decreased as the learning progresses.
- ▶ In other words, fluctuations are permitted initially, but controlled with the progress of the iterations – simulated annealing idea.
- ▶ It is possible to take best of both worlds and perform mini-batch updates – gradients are calculated for a batch of  $n$  training examples.
- ▶ Mini-Batch updates provide better exploitation of GPU and distributed computing

# Challenges in Mini-Batch SGD

- ▶ How to choose appropriate learning rate? A high value may lead to undesirable fluctuations.
- ▶ What may a better schedule for reducing the learning rate? Extremely slow schedule may lead us to a better minima, but then convergence will also be extremely slow.
- ▶ Should all the parameters have the same learning rate?

# LeNet5



- ▶ Trained using 32x32 pixel size images from MNIST (at most 20x20 characters centered in 28\*28)
- ▶ Cx are convolutional layers (C1, C3, C5)
- ▶ Sx are sumsampling (pooling) layers (S1, S4)
- ▶ One fully connected hidden layer (F6) with 84 neurons and an output layer with 10 neurons

# Convolutional Neural Networks

# High dimensional input

- ▶ Add one layer  $k$  with  $D_k$  neurons.
- ▶ Add  $(D_{k-1} + 1) \times D_k$  parameters –  $W^{(k)}$  and bias.
- ▶ Optimization become intractable.
- ▶ Images input is very high-dimensional.
- ▶ Each pixel of an image is a feature. If image is 100 by 100 pixels, then there are 10,000 features.

# Convolutional neural networks

- ▶ Reduces the number of parameters in a deep neural network with many units without losing too much in the quality of the model.
- ▶ CNNs have found applications in image and text processing.

# Learn locally

- ▶ In images, pixels that are close to one another usually represent the same type of information: sky, water, leaves, fur, bricks.
- ▶ Exception from the rule are the edges: the parts of an image where two different objects “touch” one another.
- ▶ Train the network to recognize regions of the same information as well as the edges  $\Rightarrow$  Can predict the object represented in the image.
- ▶ Split the image into square patches using a moving window approach.
- ▶ Learn multiple smaller regression models at once.
- ▶ Train small regression model for a square patch as input – learns to detect a specific kind of pattern in the input patch.
- ▶ One to detect the sky; another one for the grass, the third one for edges of a building.

# Convolution

- ▶ Input image is made of black and white pixels.
- ▶ A patch is a  $3 \times 3$ .
- ▶ Learn a regression model to detect a cross pattern in patches.

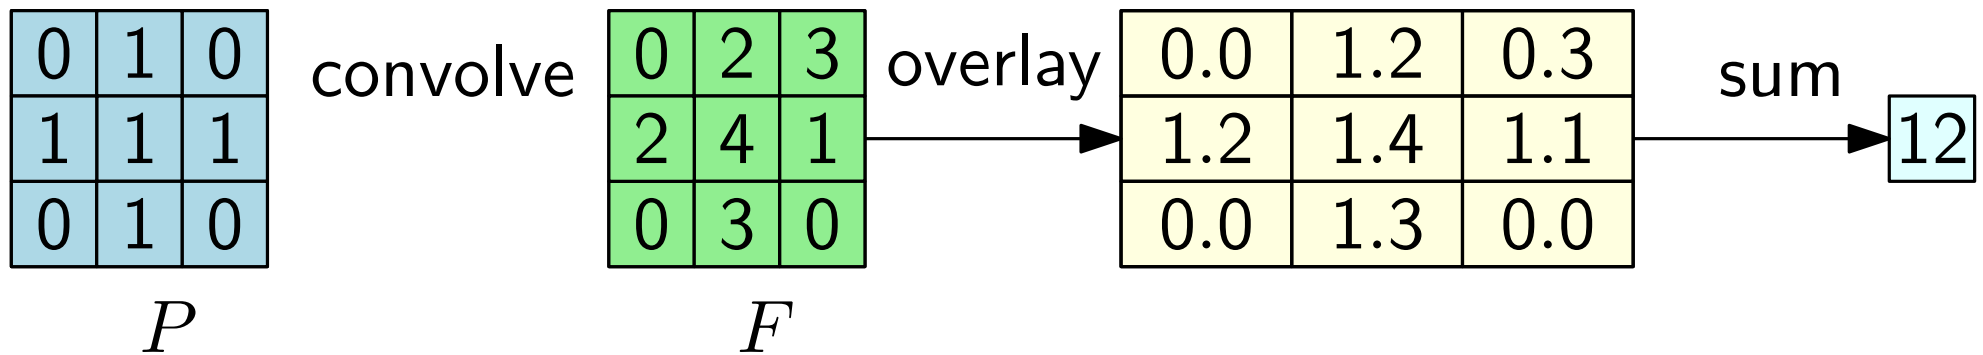
$$\mathbf{P} = \begin{bmatrix} 0 & 1 & 0 \\ 1 & 1 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$

- ▶ Learn a  $3 \times 3$  parameter matrix  $\mathbf{F}$  (**filter**): parameters corresponding to 1's will be positive and corresponding to 0's will be close to 0.

$$\mathbf{F} = \begin{bmatrix} 0 & 2 & 3 \\ 2 & 4 & 1 \\ 0 & 3 & 0 \end{bmatrix}$$



# Convolution between two matrices



$$\mathbf{P} \text{ convolution } \mathbf{F} = 12$$


- ▶ Higher the convolution value, the more similar  $\mathbf{F}$  is to  $\mathbf{P}$ .
- ▶ A bias added before applying activation function.

$$\mathbf{P} = \begin{bmatrix} 1 & 1 & 0 \\ 0 & 1 & 0 \\ 0 & 1 & 0 \end{bmatrix}, \mathbf{F} = \begin{bmatrix} 0 & 2 & 3 \\ 2 & 4 & 1 \\ 0 & 3 & 0 \end{bmatrix}$$

$$\mathbf{P} \text{ convolution } \mathbf{F} = 9$$

- ▶ Each filter slides (convolves) across the input image, left to right, top to bottom, and convolution is computed at each slide.


# A filter convolves across an image



# Convolution layer


- ▶ One convolution layer made of multiple filters.
- ▶ Filter matrices and bias are parameters learned through gradient descent.
- ▶ ReLU activationfunction in convolution layer.
- ▶ Output layer activation function?
- ▶ In convolution layer  $k$ , each filter outputs one matrix,  $D_k$  filters output  $D_k$  matrices.
- ▶ If layer  $k + 1$  is convolutional, it treats the  $D_k$  matrices as  $D_k$  image matrices, **volume of depth  $D_k$** .
- ▶ Convolution of a patch of a volume is the sum of convolutions of the corresponding patches of individual matrices of the volume.
- ▶ Input image is a volume of 3 channels: R, G, B.

# Convolution of a volume



$$(-2 \cdot 3 + 3 \cdot 1 + 5 \cdot 4 + -1 \cdot 1) + (-2 \cdot 2 + 3 \cdot -1 + 5 \cdot 0 + -1 \cdot 1) + (-2 \cdot 1 + 3 \cdot -1 + 5 \cdot 2 + -1 \cdot -1) + (-2)$$

# Stride



- Step size of the sliding window.
- Larger the slide, the smaller the output matrix.

# Padding

Conv 1

0	0	0	0	0	0
0	1	0	0	1	0
0	1	0	1	0	0
0	1	1	0	0	0
0	0	1	0	1	0
0	0	0	0	0	0

-1	2
4	-2

Filter

1
---

Bias

-1		

Conv 2

0	0	0	0	0	0
0	1	0	0	1	0
0	1	0	1	0	0
0	1	1	0	0	0
0	0	1	0	1	0
0	0	0	0	0	0

-1	2
4	-2

1
---

-1	0	

- ▶ Border around the image – how many cells wide?
- ▶ The larger the padding, the larger the output matrix.
- ▶ Helpful with larger filters – allows to better “scan” the boundaries of the image.

# Pooling with filter size 2 and stride 2

Pool 1

3	8	1	4
5	2	6	-1
-3	5	9	1
4	5	7	2

8	

Pool 3

3	8	1	4
5	2	6	-1
-3	5	9	1
4	5	7	2

8	6
5	

Pool 2

3	8	1	4
5	2	6	-1
-3	5	9	1
4	5	7	2

8	6

Pool 4

3	8	1	4
5	2	6	-1
-3	5	9	1
4	5	7	2

8	6
5	9

- ▶ Pooling applies a fixed operator, max or average – no parameters to learn. Max pooling popular.
- ▶ Hyperparameters: filter size, stride. Usually, pooling layer follows a convolution layer.
- ▶ Reduces the number of parameters, speeds up training.

# Recurrent Neural Networks




# Recurrent Neural Networks (RNN)

- ▶ Used to label, classify, or generate sequences.
- ▶ A sequence is a matrix: each row is a feature vector and the order of rows matters.
- ▶ To label a sequence is to predict a class for each feature vector in a sequence.
- ▶ To classify a sequence is to predict a class for the entire sequence.
- ▶ To generate a sequence is to output another sequence (of a possibly different length) somehow relevant to the input sequence.


# Applications

- ▶ Text processing
- ▶ Speech processing


# Recurrent neurons




Recurrent neuron unrolled in time



# Recurrent layer




# Memory cell



$$h^{(t)} = f(h^{(t-1)}, x^{(t)})$$

$$y^{(t)} = g(h^{(t-1)}, x^{(t)})$$



# RNN



Output of a RNN layer for one example.

$$\mathbf{h}^{(t)} = \phi \left( \mathbf{x}^{(t)\top} \cdot \mathbf{W}_x + \mathbf{h}^{(t-1)\top} \cdot \mathbf{W}_h + \mathbf{b} \right)$$

Output of a RNN layer for a mini-batch.


$$\mathbf{Y}^{(t)} = \phi \left( \mathbf{X}^{(t)} \cdot \mathbf{W}_x + \mathbf{Y}^{(t-1)} \cdot \mathbf{W}_h + \mathbf{b} \right)$$

$$= \phi \left( \begin{bmatrix} \mathbf{X}^{(t)} & \mathbf{Y}^{(t-1)} \end{bmatrix} \cdot \mathbf{W} + \mathbf{b} \right)$$




# Input Output sequences


one to one




one to seq




seq to one



delayed seq to seq



seq to seq



- ▶ one to one: Image classification
- ▶ one to seq : Image captioning
- ▶ seq to one: Text classification
- ▶ delayed seq to seq: Translation
- ▶ seq to seq: Time series prediction

# Difficulties

- ▶ **Propagation through time** is used to compute the parameters.
- ▶ Because of the sequential nature of the input, backpropagation has to “unfold” the network over time.
- ▶ The longer the input sequence, the deeper is the unfolded network.
- ▶ Even if our RNN has just one or two recurrent layers, both  $\tanh$  and *softmax* suffer from the vanishing gradient problem.
- ▶ **Long-term dependencies.**
- ▶ The feature vectors from the beginning of the sequence tend to be “forgotten”, because the state of each unit becomes significantly affected by the feature vectors read more recently.
- ▶ In text or speech processing, the cause-effect link between distant words in a long sentence can be lost.



# Gated RNN

- ▶ Long short-term memory (LSTM) networks.
- ▶ Networks based on the Gated Recurrent Unit (GRU).
- ▶ Units make decisions about what information to store, and when to allow reads, writes, and erasures.
- ▶ Those decisions are learned from data and implemented through the concept of gates.
- ▶ There are several architectures of gated units.
- ▶ **Minimal gated GRU** and is composed of a memory cell, and a forget gate.
- ▶ Can store information in their units for future use.
- ▶ Reading, writing, and erasure of information stored in each unit is controlled by activation functions – values in the range  $(0, 1)$ .

Tools

# Python



▶ <https://www.python.org/>

- ▶ Very high-level data structures.
- ▶ Clean syntax.
- ▶ Eco-system (comes with batteries attached!)
- ▶ De Facto programming language of ML.
- ▶ A large collection of ML packages.
- ▶ Python 2.7, Python 3

# Anaconda



▶ <https://anaconda.org/>

- ▶ Most popular Python data science platform.
- ▶ Leads open source projects like Anaconda, NumPy and SciPy that form the foundation of modern data science.

# NumPy



▶ <http://www.numpy.org/>

Fundamental package for scientific computing with Python.

- ▶ A powerful N-dimensional array object.
- ▶ Sophisticated functions.
- ▶ Tools for integrating C/C++ and Fortran code.
- ▶ Useful linear algebra, Fourier transform, and random number capabilities.

# scikit-learn



▶ <http://scikit-learn.org/stable/>

- ▶ Classification: SVM, Nearest neighbors, Random forests.
- ▶ Regression: SVR, Ridge regression, Lasso.
- ▶ Clustering: k-means, Spectral clustering.
- ▶ PCA, Feature selection, Non-negative matrix factorization.
- ▶ Model selection: Grid search, Cross validation, Metrics.
- ▶ Preprocessing: Feature extraction.

# TensorFlow



- ▶ Open source software library for numerical computation using data flow graphs.
- ▶ Nodes in the graph represent mathematical operations, while the graph edges represent the multidimensional data arrays (tensors) communicated between them.
- ▶ Deploy computation to one or more CPUs or GPUs in a desktop, server, or mobile device with a single API.
- ▶ Originally developed by Google Brain Team for machine learning and deep neural networks research.
- ▶ General enough to be applicable in a wide variety of other domains as well.

# Keras



- ▶ High-level neural networks API, written in Python and capable of running on top of TensorFlow, CNTK, or Theano.
- ▶ Allows for easy and fast prototyping.
- ▶ Supports both convolutional networks and recurrent networks, as well as combinations of the two.
- ▶ Runs seamlessly on CPU and GPU.



# PyTorch



- ▶ Open source machine learning framework based on the Torch library.
- ▶ Applications such as computer vision and natural language processing.
- ▶ Developed by Facebook's AI Research lab (FAIR).
- ▶ Although the Python interface is more polished and the primary focus of development, PyTorch also has a C++ interface.

# References

- ▶ Ian Goodfellow, Yoshua Bengio, Aaron Courville, “Deep Learning”, 2016
- ▶ Eli Stevens, Luca Antiga, Thomas Viehmann, “Deep Learning with PyTorch”, 2020

Questions?